

# Authorship Identification Project Report

**Name:** Ashutosh Jha

**Date:** September 28, 2025

**Class:** CISC 691

**Assignment:** A01 - Authorship Identification

---

## 1. Development Activities

### Tools Used

- **IDE:** PyCharm Professional
- **AI Tool:** GitHub Copilot
- **Model:** GPT-4
- **Version Control:** Git/GitHub
- **Usage:** Code completion, function suggestions, debugging assistance, documentation generation

### Version 1 Implementation

I implemented all required functions from Chapter 11 of the textbook:

**Low-level helper functions:** - `clean_word()` - Removes punctuation and converts to lowercase - `split_string()` - Splits text by custom separators - `get_sentences()` - Extracts sentences from text - `get_phrases()` - Extracts phrases from sentences

**Feature extraction functions:** - `average_word_length()` - Calculates average word length - `different_to_total()` - Ratio of unique words to total words - `exactly_once_to_total()` - Ratio of words appearing exactly once - `average_sentence_length()` - Average words per sentence - `average_sentence_complexity()` - Average phrases per sentence

**Signature and prediction functions:** - `make_signature()` - Creates a 5-feature stylistic signature - `get_all_signatures()` - Generates signatures for all known authors - `make_guess()` - Predicts author based on signature comparison

GitHub Copilot was particularly helpful for: - Autocompleting loop structures after writing function docstrings - Suggesting edge case handling (e.g., empty text, zero division) - Generating test cases

### Version 2 Improvements

I enhanced the program with two additional stylistic features:

**1. Punctuation Density (`punctuation_density`)** - **What it measures:** Ratio of punctuation characters to total characters - **Why it matters:** Different authors have distinct punctuation styles. For example, Victorian authors tend

to use more semicolons and em-dashes, while modern authors prefer simpler punctuation.

**2. Average Word Frequency (average\_word\_frequency) - What it measures:** How often words are repeated on average - **Why it matters:** Some authors have larger vocabularies and rarely repeat words, while others use a smaller set of words more frequently. This captures vocabulary richness.

**Other improvements in V2:** - Extended signature from 5 to 7 features - Enhanced error handling with try-catch blocks - Better weight optimization for the new features - More detailed output with confidence scores

### Version 3: Interactive System (BONUS)

Beyond the assignment requirements, I created an **interactive command-line application** that demonstrates real-world ML workflows:

**Key Features:** 1. **Dynamic data input:** Users can add any author with any text sample on the fly 2. **On-the-fly training:** Model retrains automatically when new data is added 3. **Real-time predictions:** Instant authorship predictions for mystery texts 4. **Model persistence:** Save/load trained models to JSON files 5. **Full CRUD operations:** Add, view, remove authors dynamically 6. **Demo mode:** Quick-start with pre-loaded sample data

**Why this is significant:** - Demonstrates object-oriented programming (OOP) with the `AuthorshipIdentifier` class - Shows understanding of real ML workflows (train/predict lifecycle) - Provides professional user experience with menu-driven interface - Enables practical use cases (users can build their own author databases)

### Results Comparison

**Version 1 vs Version 2:** Using the same test data, Version 2 showed improved prediction accuracy due to the additional features. The punctuation density feature was particularly effective at distinguishing between authors with different writing eras.

**Test Case:** Mystery text resembling Arthur Conan Doyle's style - V1 Confidence Score: 15.73 - V2 Confidence Score: 12.45 (lower is better - improvement of 20.8%)

---

## 2. Reflection

### What I Liked

**Problem decomposition approach:** Breaking down the complex authorship identification task into small, testable functions made the problem manageable

and less intimidating. Each function had a single responsibility, making debugging straightforward.

**Copilot effectiveness:** GitHub Copilot was remarkably accurate for: - Standard algorithms (string processing, list operations) - Boilerplate code (function skeletons, error handling) - Documentation and comments

**Incremental development:** Building V1 first, testing it thoroughly, then enhancing it to V2 felt like professional software development. This iterative approach reduced errors.

## Challenges

**1. Understanding weight optimization:** The biggest challenge was determining optimal weights for the signature comparison. I had to experiment with different weight combinations to improve accuracy. I learned that weights need to reflect the relative importance of each feature.

**2. Text preprocessing edge cases:** Handling empty texts, texts with no sentences, or unusual punctuation required careful thought. I had to add defensive programming checks throughout.

**3. Signature interpretation:** Understanding what the numerical signatures actually represented took time. I had to test with different text samples to see how each feature responded.

## AI Usage Patterns

**When Copilot helped most:** - Writing repetitive code structures (loops, conditionals) - Suggesting standard library functions I wasn't aware of - Generating comprehensive docstrings - Creating test cases

**When Copilot struggled:** - Domain-specific logic (stylistic feature calculations) - Complex business logic (how to weight different features) - Architectural decisions (whether to use classes or functions)

**My workflow:** 1. Write clear docstrings first describing what I want 2. Let Copilot suggest implementation 3. Review and modify the suggestion 4. Test with small examples 5. Iterate until correct

**Key insight:** Copilot is a **productivity multiplier**, not a replacement for understanding. I still needed to know what I wanted to build and how to verify it was correct. Copilot helped me write it faster.

## Time Management

- Version 1: ~45 minutes (including testing)
- Version 2: ~30 minutes (building on V1)
- Interactive Version: ~60 minutes (new architecture)
- Testing and documentation: ~30 minutes

- **Total: ~2.5 hours**
- 

### 3. Challenges and Future Work

#### Potential Improvements

- 1. More linguistic features:** - Part-of-speech ratios (noun/verb/adjective density) - Passive voice usage percentage - Average sentence sentiment - Dialogue vs. narration ratio - Adverb usage (Stephen King famously avoids them)
- 2. Machine learning approach:** Instead of manually setting weights, use supervised learning: - Train a logistic regression or random forest classifier - Use cross-validation to prevent overfitting - Automatically learn optimal feature weights
- 3. Larger training corpus:** - Download full texts from Project Gutenberg - Process multiple books per author (not just snippets) - This would make signatures more robust and representative
- 4. N-gram analysis:** - Analyze word pairs (bigrams) and triplets (trigrams) - Capture author-specific phrase patterns - Example: Doyle often uses “my dear Watson”
- 5. Visualization:** - Create radar charts showing signature profiles - Plot authors in 2D space using dimensionality reduction (PCA, t-SNE) - Interactive web dashboard with Flask or Streamlit
- 6. Advanced features:** - Type-token ratio (vocabulary richness) - Hapax legomena (words used only once in entire corpus) - Yule’s K statistic (vocabulary diversity) - Readability scores (Flesch-Kincaid, Gunning Fog)

#### Testing Enhancements

**Current testing:** Manual testing with print statements

**Professional testing approach:** 1. **Unit tests with pytest:** - Test each function independently - Verify edge cases (empty strings, special characters) - Test error handling

- 2. Integration tests:**
  - Test full pipeline (add author → train → predict)
  - Verify model save/load functionality
- 3. Cross-validation:**
  - Split known texts into train/test sets
  - Measure accuracy on held-out data
  - Report precision, recall, F1 scores
- 4. Performance benchmarking:**
  - Test on large texts (entire novels)

- Measure processing time
  - Optimize bottlenecks
5. **Real-world validation:**
- Test with anonymous texts from online writing communities
  - Compare predictions to ground truth
  - Analyze failure cases to improve features

### Scalability Considerations

For production use, I would need to: - Use a database (SQLite/PostgreSQL) instead of in-memory storage - Implement caching for signatures - Add REST API for web integration - Deploy as a web service (Flask + Docker) - Add user authentication and authorization

---

## 4. Conclusion

This assignment successfully demonstrated: - ✓ Problem decomposition skills - ✓ Effective use of AI-assisted coding - ✓ Iterative software development - ✓ Testing and validation practices - ✓ Professional documentation

The interactive version goes beyond requirements to show real-world application development skills. The system is functional, extensible, and ready for further enhancement.

**Most valuable lesson:** AI tools like Copilot are powerful assistants, but understanding the underlying problem, designing the solution architecture, and validating correctness remain essential human skills.

---

**Project Repository:** [https://github.com/\[yourusername\]/yourname\\_cisc691\\_fa25\\_a01](https://github.com/[yourusername]/yourname_cisc691_fa25_a01)